Cao Le Minh Khoa

Ho Chi Minh City, Vietnam | godminhkhoa@gmail.com | LinkedIn

Experience

AgileMark **Jul 2026 – Present**
Junior AI Engineer – AI Security & Cybersecurity *Remote, Singapore*

- Architected and built, as the **sole engineer**, an enterprise RFI and compliance intelligence platform for **banking/GRC** workflows, transforming heterogeneous requests and regulatory documents into structured requirements and evidence-backed assessments.
- Built a grounded AI pipeline spanning document parsing, LLM-based extraction, hybrid retrieval and reranking, provenance verification, structured judgments, source-backed rationales, and human review across document, spreadsheet, email, and image evidence.
- Deployed the system with **FastAPI, PostgreSQL, Qdrant, Docker, Ollama, and local Qwen models**; owned backend APIs, evaluation, deployment, and iterative redesign through direct customer demos and feedback.

EyeCode Tech **Jun 2026 – Aug 2026**
AI Researcher Intern *Ho Chi Minh City*

- Researched efficient **Vision-Language-Action (VLA)** inference for SmolVLA, implementing real forward-path token pruning, transformer layer bypass, runtime telemetry, and ablation pipelines across LeRobot/SO100 and LIBERO-10; later redesigned the approach toward trainable context-conditioned token and layer routing.
- Built reproducible efficiency evaluations spanning compute, latency, and task quality: a joint step-aware ablation reduced normalized compute from **1.00 to ∼0.45** and rollout latency from **344 ms to 317 ms**; on LIBERO-10, reducing flow-matching steps from 10 to 5 cut latency by **28.25%** (236.97 ms to 170.03 ms) while maintaining aggregate performance over 50 episodes.

Fizo **Feb 2025 – Present**
Fresher AI Engineer – Finance / Fintech *On-site, Ho Chi Minh City*

- Built and operated an AI-powered customer-service automation system integrated with **Zalo OA**, used in internal customer-service operations.
- Owned the backend and infrastructure implementation in a 2-person team, connecting Zalo OA webhooks to server-hosted **n8n** workflows for knowledge retrieval, LLM response generation, automated replies, and Redis-backed cache, queue, and session state.
- Designed a human-in-the-loop fallback that routes unsupported questions to CS staff and syncs verified responses back into the knowledge flow; deployed the service on **Linux** with **Docker/Nginx**, FastAPI, PostgreSQL, Qdrant, Redis, and Ollama.

Publications

MonoLaDN: Language-Guided DeNoising for Monocular 3D Object Detection
Contribution: VLM-based 3D reasoning generation **ACCV 2026 – Accepted**

- Owned **VLM 3D reasoning generation**, combining object-centric views with geometry-to-language mapping to produce structured **occlusion, depth, dimension, and angle** reasoning for training-time supervision.
- MonoLaDN reached **29.67 / 20.45 / 18.02 AP**$_{3D|R40}$ on KITTI test (easy/moderate/hard), improving over MonoDGP by **+3.32 / +1.73 / +2.05** at the same **42.4 ms** runtime because the VLM path is removed at inference.

Education

Ho Chi Minh City University of Technology (HCMUT), VNU-HCM *Ho Chi Minh City, Vietnam*
M.S. in Computer Science **2026 – 2028 (Expected)**
B.S. in Computer Science **2023 – 2027 (Expected)**

Awards

Zalo AI Challenge 2025 – Runner-up (Joint 2nd Place), AeroEyes Track **2025**

Technical Skills

Programming / ML: Python, SQL, PyTorch, VLMs, Vision-Language-Action models, monocular 3D detection
AI Systems / Infrastructure: RAG, hybrid retrieval, reranking, structured generation, FastAPI, Qdrant, PostgreSQL, Redis, Docker, Linux, n8n, Ollama